

Creating Knowledge out of Interlinked Data: The Integrated LOD2 Tool Stack

Sebastian Tramp, Bert Van Nuffelen, Philipp Frischmuth, and Sören Auer

Universität Leipzig, Institut für Informatik, AKSW,
Postfach 100920, D-04009 Leipzig, Germany,
`{lastname}@informatik.uni-leipzig.de`
<http://aksw.org>

Abstract. This joint tutorial of the consortium of the European IP project LOD2 - Creating Knowledge out of Interlinked Data will give an overview on the area of creating, managing and using Linked Data sources. As a prerequisite to the main tutorial part, we give an overview of the life cycle of Linked Data usage and its challenges as well as the LOD2 software stack, which is an integrated distribution of aligned tools which support the life-cycle of Linked Data from extraction, authoring/creation over enrichment, interlinking, fusing to visualization and maintenance. Based on this more theoretic explanations, we give a detailed insight into the usage of LOD2 stack both with practical and non-practical parts. More specifically, we present tools and usage scenarios for the following Linked Data life cycle tasks: extraction (Triplify / D2R), storage and querying (Openlink Virtuoso), authoring and manual revision (OntoWiki), interlinking and fusion (Silk / LIMES) and enrichment and repair (ORE). In addition to that, we present the LOD2 stack demonstrator which integrates these and other LOD2 tools for a better user experience.

1 Introduction

LOD2 is a large-scale integrating project co-funded by the European Commission within the FP7 Information and Communication Technologies Work Programme (Grant Agreement No. 257943). Commencing in September 2010, this 4-year project comprises leading Linked Open Data technology researchers, companies, and service providers (15 partners) from across 11 European countries (and one associated partner from Korea) and is coordinated by the AKSW research group at the University of Leipzig.

The goals of the LOD2 project include the development of enterprise-ready tools and methodologies for exposing and managing very large amounts of structured information on the Data Web and adaptive tools




Fig. 1. Linked Data Lifecycle

for searching, browsing, and authoring of Linked Data. These tools are distributed as the LOD2 stack. The LOD2 stack is an integrated distribution of aligned tools which support the life-cycle of Linked Data (depicted in Figure 1) from extraction, authoring/creation over enrichment, inter-linking, fusing to visualization and maintenance.

2 The Tutorial

The tutorial will give an overview to the complete LOD2 stack and a detailed insight to the integrated tools in Sec.4. For each tool we will give a talk followed by a practical part. Each practical part is guided by an overall usage scenario which covers most parts of the Linked Data life cycle. For the practical parts, we provide a virtual machine where the LOD2 stack and additional material is already installed (slides, data sources, software packages).

The intended audience covers knowledge workers from the industry which need to integrate and use Linked Data from different open and non-open data sources as well as researchers and PhD students which are new in the area of Linked Data.

This is a full day tutorial and it was never given before on other venues.

3 The Presenters

Dr. Sören Auer

Dr. Bert Van Nuffelen

Sebastian Tramp studied information science and is a research associate at AKSW / University of Leipzig since 2006. He is author of over 30 peer-reviewed publications and presented his research projects on conferences as the ESWC and the ISWC as well as multiple workshops. He is a lecturer at the University of Leipzig and the Leipzig School of Media.

4 The LOD2 Stack

The LOD2 stack comprises new and substantially extended existing tools from the LOD2 partners and third parties. The LOD2 stack is organized as a Debian package repository making the tool stack easy to install on any Debian-based system (e.g. Ubuntu). In our tutorial, we present tools and usage scenarios for the following Linked Data life cycle tasks:

4.1 Extraction

Triplify [2] provides a building block for the "semantification" of Web applications. Triplify is a small plugin for Web applications, which reveals the semantic structures encoded in relational databases by making database content available as RDF, JSON or Linked Data.

The D2R Server [3] is a tool for publishing the content of relational databases on the Semantic Web. D2R Server uses a customizable D2RQ mapping to map database content into RDF triple, and allows the RDF knowledge bases to be browsed and searched - the two main access paradigms to the Semantic Web.

4.2 Storage and Querying

Virtuoso [4] is an innovative industry standards compliant platform for native data, information, and knowledge management. It implements and supports a broad spectrum of query languages, data access interfaces, protocols, and data representation formats that includes SPARQL, RDF, RDFa and many more. The open-source edition of Virtuoso, which includes a scalable high-performance RDF Quad Store, is the basis for the LOD2 Stack's knowledge store.

4.3 Authoring and Manual Revision

OntoWiki [1] is a tool providing support for agile, distributed knowledge engineering scenarios. It facilitates the visual presentation of a knowledge base as an information map, with different views on instance data. It enables intuitive authoring of semantic content, with an inline editing mode for editing RDF content, similar to WYSIWIG for text documents. OntoWiki provides sophisticated means for navigating, visualising and authoring of RDF-based Knowledge Bases. It serves and consumes Linked Data and comprises a comprehensive middleware API for building custom Semantic Web applications.

4.4 Interlinking and Fusion

The Silk Link Discovery Framework [7] supports data publishers in interlinking RDF knowledge bases. Using the declarative Silk - Link Specification Language (Silk-LSL), developers can specify which types of RDF links should be discovered between data sources as well as which conditions data items must fulfill in order to be interlinked. These link conditions may combine various similarity metrics and can take the graph around a data item into account, which is addressed using an RDF path language.

LIMES [6] is a link discovery framework for the Web of Data. It implements time-efficient approaches for large-scale link discovery based on the characteristics of metric spaces. It is easily configurable via a web interface. It can also be downloaded as standalone tool for carrying out link discovery locally.

4.5 Enrichment and Repair

ORE [5] is a tool for repairing and enriching OWL ontologies. State-of-the-art methods in ontology debugging and supervised machine learning form the basis of ORE and are adapted or extended so as to work well in practice. ORE supports the detection of a variety of ontology modelling problems and guides the user through the process of resolving them. Furthermore, the tool allows to extend an ontology through (semi-)automatic supervised learning. A wizard-like process helps the user to resolve potential issues after axioms are added.

Acknowledgments This work was supported by a grant from the European Union's 7th Framework Programme provided for the project LOD2 (GA no. 257943).

References

1. S. Auer, S. Dietzold, and T. Riechert. OntoWiki - A Tool for Social, Semantic Collaboration. In *ISWC 2006*, volume 4273 of *LNCIS*. Springer, 2006.
2. Sören Auer, Sebastian Dietzold, Jens Lehmann, Sebastian Hellmann, and David Aumüller. Triplify: Light-weight linked data publication from relational databases. In Juan Quemada, Gonzalo León, Yoëlle S. Maarek, and Wolfgang Nejdl, editors, *Proceedings of the 18th International Conference on World Wide Web, WWW 2009, Madrid, Spain, April 20-24, 2009*, pages 621–630. ACM, 2009.
3. Christian Bizer. D2r map - a database to rdf mapping language, 2003.
4. Orri Erling and Ivan Mikhailov. Rdf support in the virtuoso dbms. In Sören Auer, Christian Bizer, Claudia Müller, and Anna V. Zhdanova, editors, *The Social Semantic Web 2007, Proceedings of the 1st Conference on Social Semantic Web (CSSW), September 26-28, 2007, Leipzig, Germany*, volume 113 of *LNI*, pages 59–68. GI, 2007.
5. Jens Lehmann and Lorenz Bühmann. Ore - a tool for repairing and enriching knowledge bases. In *Proceedings of the 9th International Semantic Web Conference (ISWC2010)*, Lecture Notes in Computer Science, Berlin / Heidelberg, 2010. Springer.
6. Axel-Cyrille Ngonga Ngomo and Sören Auer. Limes a time-efficient approach for large-scale link discovery on the web of data. In Toby Walsh, editor, *IJCAI 2011, Proceedings of the 22nd International Joint Conference on Artificial Intelligence, Barcelona, Catalonia, Spain, July 16-22, 2011*, pages 2312–2317. IJCAI/AAAI, 2011.
7. Julius Volz, Christian Bizer, Martin Gaedke, and Georgi Kobilarov. Discovering and maintaining links on the web of data. In Abraham Bernstein, David R. Karger, Tom Heath, Lee Feigenbaum, Diana Maynard, Enrico Motta, and Krishnaprasad Thirunarayan, editors, *The Semantic Web - ISWC 2009, 8th International Semantic Web Conference, ISWC 2009, Chantilly, VA, USA, October 25-29, 2009. Proceedings*, volume 5823 of *Lecture Notes in Computer Science*, pages 650–665. Springer, 2009.